

K- Nearest Neighbors on Amazon Food Reviews

```
In [20]: %matplotlib inline
import warnings
warnings.filterwarnings("ignore")

# General Packages
import os
import sqlite3
import pandas as pd
import numpy as np
import string
import re
import nltk
import datetime
import time

# Plotting Packages
import matplotlib.pyplot as plt
import seaborn as sns

# Packages for Tfidf
from sklearn.feature_extraction.text import TfidfTransformer
from sklearn.feature_extraction.text import TfidfVectorizer

# Packages for BOW (Bag of words)
from sklearn.feature_extraction.text import CountVectorizer
from sklearn.metrics import confusion_matrix
from sklearn import metrics
from sklearn.metrics import roc_curve, auc

# Packages for Text Preprocessing
from nltk.stem.porter import PorterStemmer
from nltk.corpus import stopwords
from nltk.stem.wordnet import WordNetLemmatizer

#Packages for Word2vec, Average Word2vec & Tf-Idf Weighted Word2Vec
from gensim.models import Word2Vec
from gensim.models import KeyedVectors
import pickle

#Packages for plotting Tsne plot
from sklearn.manifold import TSNE

from sklearn.cross_validation import train_test_split
from sklearn.neighbors import KNeighborsClassifier
from sklearn.metrics import accuracy_score
from sklearn.cross_validation import cross_val_score
from collections import Counter
from sklearn.metrics import accuracy_score
from sklearn import cross_validation
from sklearn.metrics import classification_report, confusion_matrix
from sklearn.model_selection import TimeSeriesSplit
from prettytable import PrettyTable
from sklearn.metrics import recall_score
from sklearn.metrics import precision_score
from sklearn.metrics import f1_score
from sklearn.metrics import roc_curve
from sklearn.model_selection import GridSearchCV
from sklearn.metrics import make_scorer, accuracy_score
from sklearn.decomposition import TruncatedSVD
```

```
In [2]: # Formaing Output using pretty table
x = PrettyTable(["Model", "Algorithm", " K ", "TPR", "TNR", "Precision", "Recall", "F1-Score", "FPR", "FNR", "PPV", "NPV", "Overall Accuracy (ACC)"])
from prettytable import MSWORD_FRIENDLY
x.set_style(MSWORD_FRIENDLY)
print(x)
```

Model	Algorithm	K	TPR	TNR	Precision	Recall	F1-Score	FPR	FNR	PPV	NPV	Overall Accuracy (ACC)
-------	-----------	---	-----	-----	-----------	--------	----------	-----	-----	-----	-----	------------------------

Preprocessing Stage: Cleansed Stop Words, Punctuations & Html tags

```
In [3]: #Connecting the Sqlite file after the Preprocessing Stage
os.chdir('/Users/sujis/Downloads/AI/Sujit')
con = sqlite3.connect('final.sqlite')
final= pd.read_sql_query(""" SELECT * FROM Reviews """, con)
```

```
In [4]: final.head(1)
```

Out[4]:		index	Id	ProductId	UserId	ProfileName	HelpfulnessNumerator	HelpfulnessDenominator	Score	Time	Summary	Text	CleanedText
	0	138706	150524	0006641040	ACITT7Di8IDDL	shari zychinski	0	0	positive	939340800	EVERY book is educational	this witty little book makes my son laugh at l...	witti littl book make son laugh loud recit car...

```
In [5]: final['Score'].value_counts()
```

```
Out[5]: positive      307061
negative      57110
Name: Score, dtype: int64
```

Considering 50K data for KNN with 'Brute' Algorithm

```
In [209]: # Taking 75K data points from the dataset
final_dataset=final.head(50000)
```

```
In [210]: final_dataset['Score'].value_counts()
```

```
Out[210]: positive      42540
negative      7460
Name: Score, dtype: int64
```

```
In [211]: #Sorting the Dataframe with Date to apply Timebased Split
final_dataset=final_dataset.sort_values(by='Time',ascending=1)
```

```
In [212]: # create design matrix X and target vector y
X = np.array(final_dataset.iloc[:, 0:12]) # end index is exclusive
Y = np.array(final_dataset['Score']) # showing you two ways of indexing a pandas df
```

```
In [213]: #Time based split
tscv = TimeSeriesSplit(n_splits=5)
print(tscv)

TimeSeriesSplit(max_train_size=None, n_splits=5)
```

```
In [214]: #Splitting the Dataset into Train set & Test set
for train_index, test_index in tscv.split(X):
    print("TRAIN:", train_index, "TEST:", test_index)
    X_tr, X_test = X[train_index], X[test_index]
    Y_tr, Y_test = Y[train_index], Y[test_index]

TRAIN: [ 0 1 2 ... 8332 8333 8334] TEST: [ 8335 8336 8337 ... 16665 16666 16667]
TRAIN: [ 0 1 2 ... 16665 16666 16667] TEST: [16668 16669 16670 ... 24998 24999 25000]
TRAIN: [ 0 1 2 ... 24998 24999 25000] TEST: [25001 25002 25003 ... 33331 33332 33333]
TRAIN: [ 0 1 2 ... 33331 33332 33333] TEST: [33334 33335 33336 ... 41664 41665 41666]
TRAIN: [ 0 1 2 ... 41664 41665 41666] TEST: [41667 41668 41669 ... 49997 49998 49999]
```

```
In [215]: # Splitting the Train dataset into Cross Validation & Train Datasets
for train_index, test_index in tscv.split(X_tr):
    print("TRAIN:", train_index, "TEST:", test_index)
    X_train, X_cv = X[train_index], X[test_index]
    Y_train, Y_cv = Y[train_index], Y[test_index]

TRAIN: [ 0 1 2 ... 6944 6945 6946] TEST: [ 6947 6948 6949 ... 13888 13889 13890]
TRAIN: [ 0 1 2 ... 13888 13889 13890] TEST: [13891 13892 13893 ... 20832 20833 20834]
TRAIN: [ 0 1 2 ... 20832 20833 20834] TEST: [20835 20836 20837 ... 27776 27777 27778]
TRAIN: [ 0 1 2 ... 27776 27777 27778] TEST: [27779 27780 27781 ... 34720 34721 34722]
TRAIN: [ 0 1 2 ... 34720 34721 34722] TEST: [34723 34724 34725 ... 41664 41665 41666]

In [216]: print(X_train.shape)
print(X_cv.shape)
print(X_test.shape)
print('-----')
print(Y_train.shape)
print(Y_cv.shape)
print(Y_test.shape)

(34723, 12)
(6944, 12)
(8333, 12)
-----
(34723,)
(6944,)
(8333,)

In [217]: # Converting X_Train, X_cv & X_test data is to Dataframe for the ease of use
X_train_data=pd.DataFrame(X_train, columns=['index', 'Id', 'ProductId', 'UserId', 'ProfileName',
      'HelpfulnessNumerator', 'HelpfulnessDenominator', 'Score', 'Time',
      'Summary', 'Text', 'CleanedText'])

X_cv_data=pd.DataFrame(X_cv, columns=['index', 'Id', 'ProductId', 'UserId', 'ProfileName',
      'HelpfulnessNumerator', 'HelpfulnessDenominator', 'Score', 'Time',
      'Summary', 'Text', 'CleanedText'])

X_test_data=pd.DataFrame(X_test, columns=['index', 'Id', 'ProductId', 'UserId', 'ProfileName',
      'HelpfulnessNumerator', 'HelpfulnessDenominator', 'Score', 'Time',
      'Summary', 'Text', 'CleanedText'])
```

Bag of Words

```
In [218]: # create the transform
vectorizer = CountVectorizer()
# tokenize and build vocab
vectorizer.fit(X_train_data['CleanedText'].values)
# Bag of Words : Train Data Set
Train_X_vector = vectorizer.transform(X_train_data['CleanedText'].values)
# summarize encoded vector
print(Train_X_vector.shape)

(34723, 23935)

In [219]: # Cross Validation Data Set
CV_X_vector = vectorizer.transform(X_cv_data['CleanedText'].values)
print(CV_X_vector.shape)

(6944, 23935)

In [220]: #Test Data Set
Test_X_vector = vectorizer.transform(X_test_data['CleanedText'].values)
print(Test_X_vector.shape)

(8333, 23935)

In [221]: print(Train_X_vector.shape)
print(CV_X_vector.shape)
print(Test_X_vector.shape)
print('-----')
print(Y_train.shape)
print(Y_cv.shape)
print(Y_test.shape)

(34723, 23935)
(6944, 23935)
(8333, 23935)
-----
(34723,)
(6944,)
(8333,)

In [222]: start_time_code = time.time()
for i in range(1,30,2):
    # instantiate learning model (k = 30)
    knn = KNeighborsClassifier(n_neighbors=i, algorithm="brute")

    # fitting the model on crossvalidation train
    knn.fit(Train_X_vector, Y_train)

    # predict the response on the crossvalidation train
    pred = knn.predict(CV_X_vector)

    # evaluate CV accuracy
    bag_acc = accuracy_score(Y_cv, pred, normalize=True) * float(100)
    print('\nCV accuracy for k = %d is %0.1f%%' % (i, bag_acc))
end_time_code = time.time()
print ("Running Time for code execution " + str(end_time_code - start_time_code) + " secs")

CV accuracy for k = 1 is 77.0%

CV accuracy for k = 3 is 81.8%

CV accuracy for k = 5 is 82.7%

CV accuracy for k = 7 is 82.7%

CV accuracy for k = 9 is 82.7%

CV accuracy for k = 11 is 82.7%

CV accuracy for k = 13 is 82.7%

CV accuracy for k = 15 is 82.7%

CV accuracy for k = 17 is 82.6%

CV accuracy for k = 19 is 82.5%

CV accuracy for k = 21 is 82.5%

CV accuracy for k = 23 is 82.5%

CV accuracy for k = 25 is 82.5%

CV accuracy for k = 27 is 82.5%

CV accuracy for k = 29 is 82.6%
Running Time for code execution 231.88899397850037 secs

In [223]: knn = KNeighborsClassifier(15,algorithm="brute")
knn.fit(Train_X_vector,Y_train)
pred = knn.predict(Test_X_vector)
bag_test_acc = accuracy_score(Y_test, pred, normalize=True) * float(100)
print('\n***Test accuracy for k = 15 is %0.2f%%' % (bag_test_acc))

***Test accuracy for k = 15 is 81.70%

In [224]: Conf_matrix=confusion_matrix(Y_test, pred)
```

```
In [225]: Conf_matrix

Out[225]: array([[ 29, 1520],
 [ 5, 6779]])

In [226]: # Function for plotting Confusion Matrix
# Using this snipped of code copied from kaggle
import numpy as np

def plot_confusion_matrix(cm,
                          target_names,
                          title='Confusion matrix',
                          cmap=None,
                          normalize=True):

    import matplotlib.pyplot as plt
    import numpy as np
    import itertools

    accuracy = np.trace(cm) / float(np.sum(cm))
    misclass = 1 - accuracy

    if cmap is None:
        cmap = plt.get_cmap('Blues')

    plt.figure(figsize=(8, 6))
    plt.imshow(cm, interpolation='nearest', cmap=cmap)
    plt.title(title)
    plt.colorbar()


    if target_names is not None:
        tick_marks = np.arange(len(target_names))
        plt.xticks(tick_marks, target_names, rotation=45)
        plt.yticks(tick_marks, target_names)

    if normalize:
        cm = cm.astype('float') / cm.sum(axis=1)[:, np.newaxis]

    thresh = cm.max() / 1.5 if normalize else cm.max() / 2
    for i, j in itertools.product(range(cm.shape[0]), range(cm.shape[1])):
        if normalize:
            plt.text(j, i, "{0.4f}".format(cm[i, j]),
                     horizontalalignment="center",
                     color="white" if cm[i, j] > thresh else "black")
        else:
            plt.text(j, i, "{:,}".format(cm[i, j]),
                     horizontalalignment="center",
                     color="white" if cm[i, j] > thresh else "black")

    plt.tight_layout()
    plt.ylabel('True label')
    plt.xlabel('Predicted label\naccuracy={0.4f}; misclass={0.4f}'.format(accuracy, misclass))
    plt.show()

In [227]: plot_confusion_matrix(cm            = Conf_matrix,
                               normalize     = False,
                               target_names  = ['Negative', 'Positive'],
                               title         = "Confusion Matrix")
```



```
In [228]: Conf_matrix

Out[228]: array([[ 29, 1520],
 [ 5, 6779]])

In [229]: TN,FP,FN,TP = Conf_matrix.ravel()
#confusion_matrix(Y_test, pred)

In [230]: # Sensitivity, hit rate, recall, or true positive rate
TPR = '{0:.2%}'.format( TP/(TP+FN))
# Specificity or true negative rate
TNR = '{0:.2%}'.format( TN/(TN+FP))
# Fall out or false positive rate
FPR = '{0:.2%}'.format( FP/(FP+TN))
# False negative rate
FNR = '{0:.2%}'.format( FN/(FN+TP))
# Precision or positive predictive value
PPV = '{0:.2%}'.format( TP/(TP+FP))
# Negative predictive value
NPV = '{0:.2%}'.format( TN/(TN+FN))
# Overall accuracy
ACC = '{0:.2%}'.format( (TP+TN)/(TP+FP+FN+TN))

In [231]: Recall=recall_score(Y_test, pred, average='micro')

In [232]: Precision=precision_score(Y_test, pred, average='micro')

In [233]: F1_Score=f1_score(Y_test, pred, average='weighted')

In [234]: pd.crosstab(Y_test, pred, rownames=[ 'True' ], colnames=[ 'Predicted' ], margins=True)

Out[234]:
Predicted negative positive All
True
negative 29 1520 1549
positive 5 6779 6784
All 34 8299 8333

In [235]: x.add_row(["Bag of Words", "Brute", 15, TPR, TNR, '{0:.2f}'.format(Precision), '{0:.2f}'.format(Recall), '{0:.2f}'.format(F1_Score), FPR, FNR, PPV, NPV, ACC])

In [236]: print(x)
```

	Model	Algorithm	K	TPR	TNR	Precision	Recall	F1-Score	FPR	FNR	PPV	NPV	Overall Accuracy (ACC)
	Bag of Words	Kd_tree	17	99.07%	7.29%	0.82	0.82	0.76	92.71%	0.93%	82.33%	64.15%	81.94%
	TF - IDF	Kd_tree	13	99.09%	10.93%	0.83	0.83	0.77	89.07%	0.91%	82.90%	73.38%	82.64%
	Average W2V	Kd_tree	11	98.25%	18.11%	0.83	0.83	0.79	81.89%	1.75%	83.95%	70.42%	83.30%
	TF - IDF Weighted W2V	Kd_tree	23	99.34%	10.93%	0.83	0.83	0.77	89.07%	0.66%	82.94%	79.07%	82.84%
	Bag of Words	Brute	15	99.93%	1.87%	0.82	0.82	0.74	98.13%	0.07%	81.68%	85.29%	81.70%

```
In [237]: print(classification_report(Y_test, pred))
```

	precision	recall	f1-score	support
negative	0.85	0.02	0.04	1549
positive	0.82	1.00	0.90	6784
avg / total	0.82	0.82	0.74	8333

TF - IDF

```
In [238]: tf_transformer = TfidfVectorizer(ngram_range=(1,2))
TF_vector=tf_transformer.fit(X_train_data['CleanedText'].values)
TF_train_Vector = TF_vector.transform(X_train_data['CleanedText'].values)
TF_train_Vector.shape
```

Out[238]: (34723, 598035)

```
In [239]: TF_cv_Vector = TF_vector.transform(X_cv_data['CleanedText'].values)
TF_cv_Vector.shape
```

Out[239]: (6944, 598035)

```
In [240]: TF_test_Vector = TF_vector.transform(X_test_data['CleanedText'].values)
TF_test_Vector.shape
```

Out[240]: (8333, 598035)

```
In [241]: print(TF_train_Vector.shape)
print(TF_cv_Vector.shape)
print(TF_test_Vector.shape)
print('-----')
print(Y_train.shape)
print(Y_cv.shape)
print(Y_test.shape)
```

(34723, 598035)
(6944, 598035)
(8333, 598035)
-----
(34723,)
(6944,)
(8333,)

```
In [242]: start_time_code = time.time()
for i in range(1,30,2):
    # instantiate learning model (k = 30)
    knn = KNeighborsClassifier(n_neighbors=i, algorithm="brute")

    # fitting the model on crossvalidation train
    knn.fit(TF_train_Vector, Y_train)

    # predict the response on the crossvalidation train
    pred = knn.predict(TF_cv_Vector)

    # evaluate CV accuracy
    tf_acc = accuracy_score(Y_cv, pred, normalize=True) * float(100)
    print('\nCV accuracy for k = %d is %0.2f%%' % (i, tf_acc))
end_time_code = time.time()
print ("Running Time for code execution " + str(end_time_code - start_time_code) + " secs")
```

CV accuracy for k = 1 is 80.04%

CV accuracy for k = 3 is 83.19%

CV accuracy for k = 5 is 83.60%

CV accuracy for k = 7 is 83.81%

CV accuracy for k = 9 is 83.68%

CV accuracy for k = 11 is 83.80%

CV accuracy for k = 13 is 83.74%

CV accuracy for k = 15 is 83.60%

CV accuracy for k = 17 is 83.50%

CV accuracy for k = 19 is 83.40%

CV accuracy for k = 21 is 83.40%

CV accuracy for k = 23 is 83.35%

CV accuracy for k = 25 is 83.31%

CV accuracy for k = 27 is 83.35%

CV accuracy for k = 29 is 83.37%
Running Time for code execution 213.33067083358765 secs

```
In [243]: knn = KNeighborsClassifier(11,algorithm="brute")
knn.fit(TF_train_Vector,Y_train)
pred = knn.predict(TF_test_Vector)
tf_test_acc = accuracy_score(Y_test, pred, normalize=True) * float(100)
print('\n***Test accuracy for k = 11 is %0.2f%%' % (tf_test_acc))


***Test accuracy for k = 11 is 82.91%
```

```
In [244]: Conf_matrix=confusion_matrix(Y_test, pred)
```

```
In [245]: Conf_matrix
```

Out[245]: array([[ 190, 1359],
[ 65, 6719]])

```
In [246]: plot_confusion_matrix(cm = Conf_matrix,
normalize = False,
target_names = ['Negative', 'Positive'],
title = "Confusion Matrix")
```



```
In [247]: TN,FP,FN,TP = Conf_matrix.ravel()
#confusion_matrix(Y_test, pred)
```

```
In [248]: # Sensitivity, hit rate, recall, or true positive rate
TPR = '{0:.2%}'.format( TP/(TP+FN))
# Specificity or true negative rate
TNR = '{0:.2%}'.format( TN/(TN+FP))
# Fall out or false positive rate
FPR = '{0:.2%}'.format( FP/(FP+TN))
# False negative rate
FNR = '{0:.2%}'.format( FN/(TP+FN))
# Precision or positive predictive value
PPV = '{0:.2%}'.format( TP/(TP+FP))
# Negative predictive value
NPV = '{0:.2%}'.format( TN/(TN+FN))
# Overall accuracy
ACC = '{0:.2%}'.format( (TP+TN)/(TP+FP+FN+TN))
```

```
In [249]: Recall=recall_score(Y_test, pred, average='micro')
```

```
In [250]: Precision=precision_score(Y_test, pred, average='micro')
```

```
In [251]: F1_Score=f1_score(Y_test, pred, average='weighted')
```

```
In [252]: pd.crosstab(Y_test, pred, rownames=['True'], colnames=['Predicted'], margins=True)
```

Out[252]:

	Predicted negative	positive	All
True			
negative	190	1359	1549
positive	65	6719	6784
All	255	8078	8333

```
In [253]: x.add_row(["TF - IDF", "Brute", 11, TPR, TNR, '{0:.2f}'.format(Precision), '{0:.2f}'.format(Recall), '{0:.2f}'.format(F1_Score), FPR, FNR, PPV, NPV, ACC])
```

```
In [254]: print(classification_report(Y_test, pred))
```

	precision	recall	f1-score	support
negative	0.75	0.12	0.21	1549
positive	0.83	0.99	0.90	6784
avg / total	0.82	0.83	0.78	8333

```
In [255]: print(x)
```

	Model	Algorithm	K	TPR	TNR	Precision	Recall	F1-Score	FPR	FNR	PPV	NPV	Overall Accuracy (ACC)
	Bag of Words	Kd_tree	17	99.07%	7.29%	0.82	0.82	0.76	92.71%	0.93%	82.33%	64.15%	81.94%
	TF - IDF	Kd_tree	13	99.09%	10.93%	0.83	0.83	0.77	89.07%	0.91%	82.90%	73.38%	82.64%
	Average W2V	Kd_tree	11	98.25%	18.11%	0.83	0.83	0.79	81.89%	1.75%	83.95%	70.42%	83.30%
	TF - IDF Weighted W2V	Kd_tree	23	99.34%	10.93%	0.83	0.83	0.77	89.07%	0.66%	82.94%	79.07%	82.84%
	Bag of Words	Brute	15	99.93%	1.87%	0.82	0.82	0.74	98.13%	0.07%	81.68%	85.29%	81.70%
	TF - IDF	Brute	11	99.04%	12.27%	0.83	0.83	0.78	87.73%	0.96%	83.18%	74.51%	82.91%

Average W2V

```
In [256]: # Train your own Word2Vec model using your own text corpus for Training Dataset
i=0
list_of_sent=[]
for sent in X_train_data['CleanedText'].values:
    list_of_sent.append(sent.split())

# min_count = 5 considers only words that occurred atleast 5 times
Trained_model=Word2Vec(list_of_sent,min_count=5,size=50, workers=4)

w2v_words = list(Trained_model.wv.vocab)
print("number of words that occurred minimum 5 times ",len(w2v_words))
print("sample words ", w2v_words[0:50])

number of words that occurred minimum 5 times 8321
sample words  ['littl', 'book', 'make', 'son', 'laugh', 'loud', 'car', 'drive', 'along', 'alway', 'sing', 'refrain', 'hes', 'learn', 'india', 'love', 'new', 'wo
rd', 'introduc', 'silli', 'classic', 'will', 'bet', 'still', 'abl', 'memori', 'colleg', 'rememb', 'see', 'show', 'air', 'televis', 'year', 'ago', 'child', 'sist
er', 'later', 'bought', 'day', 'thirti', 'someth', 'use', 'seri', 'song', 'student', 'teach', 'preschool', 'turn', 'whole', 'school']
```

```
In [257]: # average Word2Vec
# compute average word2vec for each review.
start_time_code = time.time()
Trained_vectors = []; # the avg-w2v for each sentence/review is stored in this list
for sent in list_of_sent: # for each review/sentence
    sent_vec = np.zeros(50) # as word vectors are of zero length
    cnt_words =0; # num of words with a valid vector in the sentence/review
    for word in sent: # for each word in a review/sentence
        if word in w2v_words:
            vec = Trained_model.wv[word]
            sent_vec += vec
            cnt_words += 1
    if cnt_words != 0:
        sent_vec /= cnt_words
    Trained_vectors.append(sent_vec)
print(len(Trained_vectors))
end_time_code = time.time()
print ("Running Time for code execution " + str(end_time_code - start_time_code) + " secs")

34723
Running Time for code execution 62.595484018325806 secs
```

```
In [258]: Train_X=Trained_vectors
```

```
In [260]: # Train your own Word2Vec model using your own text corpus
i=0
list_of_sent_cv=[]
for sent in X_cv_data['CleanedText'].values:
    list_of_sent_cv.append(sent.split())
```

```
In [261]: # average Word2Vec
# compute average word2vec for each review.
start_time_code = time.time()
Cv_vectors = []; # the avg-w2v for each sentence/review is stored in this list
for sent in list_of_sent_cv: # for each review/sentence
    sent_vec = np.zeros(50) # as word vectors are of zero length
    cnt_words =0; # num of words with a valid vector in the sentence/review
    for word in sent: # for each word in a review/sentence
        if word in w2v_words:
            vec = Trained_model.wv[word]
            sent_vec += vec
            cnt_words += 1
    if cnt_words != 0:
        sent_vec /= cnt_words
    Cv_vectors.append(sent_vec)
print(len(Cv_vectors))
end_time_code = time.time()
print ("Running Time for code execution " + str(end_time_code - start_time_code) + " secs")

6944
Running Time for code execution 11.396867990493774 secs
```

```
In [262]: CV_X=Cv_vectors
```

```
In [264]: # Train your own Word2Vec model using your own text corpus
i=0
list_of_sent_test=[]
for sent in X_test_data['CleanedText'].values:
    list_of_sent_test.append(sent.split())
```

```
In [265]: # average Word2Vec
# compute average word2vec for each review.
start_time_code = time.time()
test_vectors = []; # the avg-w2v for each sentence/review is stored in this list
for sent in list_of_sent_test: # for each review/sentence
    sent_vec = np.zeros(50) # as word vectors are of zero length
    cnt_words =0; # num of words with a valid vector in the sentence/review
    for word in sent: # for each word in a review/sentence
        if word in w2v_words:
            vec = Trained_model.wv[word]
            sent_vec += vec
            cnt_words += 1
    if cnt_words != 0:
        sent_vec /= cnt_words
    test_vectors.append(sent_vec)
print(len(test_vectors))
end_time_code = time.time()
print ("Running Time for code execution " + str(end_time_code - start_time_code) + " secs")
```

8333
Running Time for code execution 14.642790079116821 secs

```
In [266]: Test_X=test_vectors
```

```
In [267]: print (len(Train_X))
print(len(CV_X))
print(len(Test_X))
print('-----')
print (Y_train.shape)
print(Y_cv.shape)
print(Y_test.shape)
```

34723
6944
8333
-----
(34723,)
(6944,)
(8333,)

```
In [268]: start_time_code = time.time()
for i in range(1,30,2):
    # instantiate learning model (k = 30)
    knn = KNeighborsClassifier(n_neighbors=i, algorithm="brute")

    # fitting the model on crossvalidation train
    knn.fit(Train_X, Y_train)

    # predict the response on the crossvalidation train
    pred = knn.predict(CV_X)

    # evaluate CV accuracy
    train_acc = accuracy_score(Y_cv, pred, normalize=True) * float(100)
    print('\nCV accuracy for k = %d is %0.2f%%' % (i, train_acc))
end_time_code = time.time()
print ("Running Time for code execution " + str(end_time_code - start_time_code) + " secs")
```

CV accuracy for k = 1 is 80.17%

CV accuracy for k = 3 is 83.45%

CV accuracy for k = 5 is 84.12%

CV accuracy for k = 7 is 84.48%

CV accuracy for k = 9 is 84.59%

CV accuracy for k = 11 is 84.53%

CV accuracy for k = 13 is 84.48%

CV accuracy for k = 15 is 84.48%

CV accuracy for k = 17 is 84.56%

CV accuracy for k = 19 is 84.52%

CV accuracy for k = 21 is 84.45%

CV accuracy for k = 23 is 84.36%

CV accuracy for k = 25 is 84.38%

CV accuracy for k = 27 is 84.25%

CV accuracy for k = 29 is 84.17%
Running Time for code execution 105.32548975944519 secs

```
In [269]: # Since the train accuracy for k is the same from k=3 to k=29. so the smallest k at maximum accuracy is the best i.e. k=3
knn = KNeighborsClassifier(17,algorithm="brute")
knn.fit(Train_X,Y_train)
pred = knn.predict(Test_X)
test_acc = accuracy_score(Y_test, pred, normalize=True) * float(100)
print('\n****Test accuracy for k = 17 is %0.2f%%' % (test_acc))
```


\*\*\*\*Test accuracy for k = 17 is 83.96%

```
In [270]: Conf_matrix=confusion_matrix(Y_test, pred)
```

```
In [271]: Conf_matrix
```

Out[271]: array([[ 323, 1226],
[ 111, 6673]])

```
In [272]: plot_confusion_matrix(cm          = Conf_matrix,
                             normalize    = False,
                             target_names = ['Negative', 'Positive'],
                             title       = "Confusion Matrix")
```



```
In [273]: TN,FP,FN,TP = Conf_matrix.ravel()
#confusion_matrix(Y_test, pred)
```

```
In [274]: # Sensitivity, hit rate, recall, or true positive rate
TPR = '{0:.2%}'.format( TP/(TP+FN))
# Specificity or true negative rate
TNR = '{0:.2%}'.format( TN/(TN+FP))
# Fall out or false positive rate
FPR = '{0:.2%}'.format( FP/(FP+TN))
# False negative rate
FNR = '{0:.2%}'.format( FN/(TP+FN))
# Precision or positive predictive value
PPV = '{0:.2%}'.format( TP/(TP+FP))
# Negative predictive value
NPV = '{0:.2%}'.format( TN/(TN+FN))
# Overall accuracy
ACC = '{0:.2%}'.format( (TP+TN)/(TP+FP+FN+TN))

In [275]: Recall=recall_score(Y_test, pred, average='micro')

In [276]: Precision=precision_score(Y_test, pred, average='micro')

In [277]: F1_Score=f1_score(Y_test, pred, average='weighted')

In [278]: pd.crosstab(Y_test, pred, rownames=['True'], colnames=['Predicted'], margins=True)

Out[278]:
      Predicted negative positive  All
True
negative      323      1226  1549
positive       111      6673  6784
All            434      7899  8333

In [279]: x.add_row([" Average W2V", "Brute", 17,TPR,TNR, '{0:.2f}'.format(Precision), '{0:.2f}'.format(Recall), '{0:.2f}'.format(F1_Score),FPR,FNR,PPV,NPV,ACC])

In [280]: print(classification_report(Y_test, pred))

              precision    recall  f1-score   support

 negative           0.74         0.21         0.33         1549
 positive           0.84         0.98         0.91         6784

 avg / total           0.83         0.84         0.80         8333

In [281]: print(x)

|      Model      | Algorithm | K |  TPR  |  TNR  | Precision | Recall | F1-Score | FPR | FNR | PPV | NPV | Overall Accuracy (ACC) |
| Bag of Words   | kd tree  | 17 | 99.07% | 7.29% | 0.82      | 0.82   | 0.76     | 92.71% | 0.93% | 82.33% | 64.15% | 81.94% |
| TF - IDF      | Kd tree  | 13 | 99.09% | 10.93% | 0.83      | 0.83   | 0.77     | 89.07% | 0.91% | 82.90% | 73.38% | 82.64% |
| Average W2V    | Kd tree  | 11 | 98.25% | 18.11% | 0.83      | 0.83   | 0.79     | 81.89% | 1.75% | 83.95% | 70.42% | 83.30% |
| TF - IDF Weighted W2V | Kd tree  | 23 | 99.34% | 10.93% | 0.83      | 0.83   | 0.77     | 89.07% | 0.66% | 82.94% | 79.07% | 82.84% |
| Bag of Words   | Brute    | 15 | 99.93% | 1.87%  | 0.82      | 0.82   | 0.74     | 98.13% | 0.07% | 81.68% | 85.29% | 81.70% |
| TF - IDF      | Brute    | 11 | 99.04% | 12.27% | 0.83      | 0.83   | 0.78     | 87.73% | 0.96% | 83.18% | 74.51% | 82.91% |
| Average W2V    | Brute    | 17 | 98.36% | 20.85% | 0.84      | 0.84   | 0.80     | 79.15% | 1.64% | 84.48% | 74.42% | 83.96% |
```

TF-IDF weighted Word2Vec

```
In [26]: # TF-IDF weighted Word2Vec
tfidf_feat = TF_vector.get_feature_names() # tfidf words/col-names

In [87]: # final_tf_idf is the sparse matrix with row= sentence, col=word and cell_val = tfidf
start_time_code = time.time()
tfidf_sent_vectors = []; # the tfidf-w2v for each sentence/review is stored in this list
row=0;
for sent in list_of_sent: # for each review/sentence
    sent_vec = np.zeros(50) # as word vectors are of zero length
    weight_sum=0; # num of words with a valid vector in the sentence/review
    for word in sent: # for each word in a review/sentence
        if word in w2v_words:
            vec = Trained_model.wv[word]
            # obtain the tf_idfidf of a word in a sentence/review
            tf_idf = TF_train_Vector[row, tfidf_feat.index(word)]
            sent_vec += (vec * tf_idf)
            weight_sum += tf_idf
        if weight_sum != 0:
            sent_vec /= weight_sum
            tfidf_sent_vectors.append(sent_vec)
    row += 1
end_time_code = time.time()
print ("Running Time for code execution " + str(end_time_code - start_time_code) + " secs")

Running Time for code execution 28771.18012213707 secs

In [88]: Train_X=tfidf_sent_vectors

In [89]: df = pd.DataFrame.from_records(Train_X)
df.to_csv('TF-IDF_W2V_Train.csv',index=None)

In [282]: Train_X = pd.read_csv("TF-IDF_W2V_Train.csv")

In [116]: len(Train_X)

Out[116]: 34723

In [92]: # final_tf_idf is the sparse matrix with row= sentence, col=word and cell_val = tfidf
start_time_code = time.time()
tfidf_sent_vectors = []; # the tfidf-w2v for each sentence/review is stored in this list
row=0;
for sent in list_of_sent_cv: # for each review/sentence
    sent_vec = np.zeros(50) # as word vectors are of zero length
    weight_sum=0; # num of words with a valid vector in the sentence/review
    for word in sent: # for each word in a review/sentence
        if word in w2v_words:
            vec = Trained_model.wv[word]
            # obtain the tf_idfidf of a word in a sentence/review
            tf_idf = TF_cv_Vector[row, tfidf_feat.index(word)]
            sent_vec += (vec * tf_idf)
            weight_sum += tf_idf
        if weight_sum != 0:
            sent_vec /= weight_sum
            tfidf_sent_vectors.append(sent_vec)
    row += 1
end_time_code = time.time()
print ("Running Time for code execution " + str(end_time_code - start_time_code) + " secs")

Running Time for code execution 5836.9589948654175 secs

In [93]: CV_X=tfidf_sent_vectors

In [94]: df = pd.DataFrame.from_records(CV_X)
df.to_csv('TF-IDF_W2V_CV.csv',index=None)

In [283]: CV_X = pd.read_csv("TF-IDF_W2V_CV.csv")
```

```
In [27]: # final_tf_idf is the sparse matrix with row= sentence, col=word and cell_val = tfidf
start_time_code = time.time()
tfidf_sent_vectors = []; # the tfidf-w2v for each sentence/review is stored in this list
row=0;
for sent in list_of_sent_test: # for each review/sentence
    sent_vec = np.zeros(50) # as word vectors are of zero length
    weight_sum =0; # num of words with a valid vector in the sentence/review
    for word in sent: # for each word in a review/sentence
        if word in w2v_words:
            vec = Trained_model.wv[word]
            # obtain the tf_idfidf of a word in a sentence/review
            tf_idf = TF_test_Vector[row, tfidf_feat.index(word)]
            sent_vec += (vec * tf_idf)
            weight_sum += tf_idf
        if weight_sum != 0:
            sent_vec /= weight_sum
            tfidf_sent_vectors.append(sent_vec)
            row += 1
end_time_code = time.time()
print ("Running Time for code execution " + str(end_time_code - start_time_code) + " secs")
```

Running Time for code execution 7043.685368061066 secs

```
In [28]: Test_X=tfidf_sent_vectors
```

```
In [29]: df = pd.DataFrame.from_records(Test_X)
df.to_csv('TF-IDF_W2V_test.csv',index=None)
```

```
In [284]: Test_X = pd.read_csv("TF-IDF_W2V_test.csv")
```

```
In [285]: start_time_code = time.time()
for i in range(1,30,2):
    # instantiate learning model (k = 30)
    knn = KNeighborsClassifier(n_neighbors=i, algorithm="brute")

    # fitting the model on crossvalidation train
    knn.fit(Train_X, Y_train)

    # predict the response on the crossvalidation train
    pred = knn.predict(CV_X)

    # evaluate CV accuracy
    train_acc = accuracy_score(Y_cv, pred, normalize=True) * float(100)
    print('\nCV accuracy for k = %d is %0.2f%%' % (i, train_acc))
end_time_code = time.time()
print ("Running Time for code execution " + str(end_time_code - start_time_code) + " secs")
```

CV accuracy for k = 1 is 80.11%

CV accuracy for k = 3 is 82.75%

CV accuracy for k = 5 is 83.55%

CV accuracy for k = 7 is 83.73%

CV accuracy for k = 9 is 83.94%

CV accuracy for k = 11 is 83.71%

CV accuracy for k = 13 is 83.67%

CV accuracy for k = 15 is 83.50%

CV accuracy for k = 17 is 83.53%

CV accuracy for k = 19 is 83.44%

CV accuracy for k = 21 is 83.48%

CV accuracy for k = 23 is 83.38%

CV accuracy for k = 25 is 83.45%

CV accuracy for k = 27 is 83.38%

CV accuracy for k = 29 is 83.41%

Running Time for code execution 100.78603434562683 secs

```
In [286]: # Since the train accuracy for k is the same from k=3 to k=29. so the smallest k at maximum accuracy is the best i.e. k=3
knn = KNeighborsClassifier(9,algorithm="brute")
knn.fit(Train_X,Y_train)
pred = knn.predict(Test_X)
test_acc = accuracy_score(Y_test, pred, normalize=True) * float(100)
print('\n****Test accuracy for k = 9 is %0.2f%%' % (test_acc))


****Test accuracy for k = 9 is 80.88%
```

```
In [287]: Conf_matrix=confusion_matrix(Y_test, pred)
```

```
In [288]: Conf_matrix
```

```
Out[288]: array([[ 32, 1517],
 [ 76, 6708]])
```

```
In [289]: plot_confusion_matrix(cm          = Conf_matrix,
                             normalize   = False,
                             target_names = ['Negative', 'Positive'],
                             title       = "Confusion Matrix")
```



```
In [290]: TN,FP,FN,TP = Conf_matrix.ravel()
#confusion_matrix(Y_test, pred)
```

```
In [291]: # Sensitivity, hit rate, recall, or true positive rate
TPR = '{0:.2%}'.format( TP/(TP+FN))
# Specificity or true negative rate
TNR = '{0:.2%}'.format( TN/(TN+FP))
# Fall out or false positive rate
FPR = '{0:.2%}'.format( FP/(FP+TN))
# False negative rate
FNR = '{0:.2%}'.format( FN/(TP+FN))
# Precision or positive predictive value
PPV = '{0:.2%}'.format( TP/(TP+FP))
# Negative predictive value
NPV = '{0:.2%}'.format( TN/(TN+FN))
# Overall accuracy
ACC = '{0:.2%}'.format( (TP+TN)/(TP+FP+FN+TN))
```



```
In [292]: Recall=recall_score(Y_test, pred, average='micro')

In [293]: Precision=precision_score(Y_test, pred, average='micro')

In [294]: F1_Score=f1_score(Y_test, pred, average='weighted')

In [295]: pd.crosstab(Y_test, pred, rownames=['True'], colnames=['Predicted'], margins=True)

Out[295]:
      Predicted  negative  positive   All
      True
negative      32      1517   1549
positive      76      6708   6784
      All      108      8225   8333

In [296]: x.add_row([" TF-IDF Weighted W2V ", "Brute", 17, TPR, TNR, '{0:.2f}'.format(Precision), '{0:.2f}'.format(Recall), '{0:.2f}'.format(F1_Score), FPR, FNR, PPV, NPV, ACC])

In [297]: print(classification_report(Y_test, pred))

              precision      recall  f1-score   support

negative      0.30      0.02      0.04      1549
positive      0.82      0.99      0.89      6784

avg / total      0.72      0.81      0.73      8333

In [298]: print(x)

      Model      Algorithm  K      TPR      TNR      Precision      Recall      F1-Score      FPR      FNR      PPV      NPV      Overall Accuracy (ACC)
      Bag of Words      kd_tree      17      99.07%      7.29%      0.82      0.82      0.76      92.71%      0.93%      82.33%      64.15%      81.94%
      TF - IDF      kd_tree      13      99.09%      10.93%      0.83      0.83      0.77      89.07%      0.91%      82.90%      73.38%      82.64%
      Average W2V      kd_tree      11      98.25%      18.11%      0.83      0.83      0.79      81.89%      1.75%      83.95%      70.42%      83.30%
      TF - IDF Weighted W2V      kd_tree      23      99.34%      10.93%      0.83      0.83      0.77      89.07%      0.66%      82.94%      79.07%      82.84%
      Bag of Words      Brute      15      99.93%      1.87%      0.82      0.82      0.74      98.13%      0.07%      81.68%      85.29%      81.70%
      TF - IDF      Brute      11      99.04%      12.27%      0.83      0.83      0.78      87.73%      0.96%      83.18%      74.51%      82.91%
      Average W2V      Brute      17      98.36%      20.85%      0.84      0.84      0.80      79.15%      1.64%      84.48%      74.42%      83.96%
      TF-IDF Weighted W2V      Brute      17      98.88%      2.07%      0.81      0.81      0.73      97.93%      1.12%      81.56%      29.63%      80.88%
```

Considering 30K data for KNN with 'kd\_tree' Algorithm

```
In [6]: # Taking 75K data points from the dataset
final_dataset=final.head(30000)

In [7]: final_dataset['Score'].value_counts()

Out[7]: positive      25494
negative      4506
Name: Score, dtype: int64

In [8]: #Sorting the Dataframe with Date to apply Timebased Split
final_dataset=final_dataset.sort_values(by='Time',ascending=1)

In [9]: # create design matrix X and target vector y
X = np.array(final_dataset.iloc[:, 0:12]) # end index is exclusive
Y = np.array(final_dataset['Score']) # showing you two ways of indexing a pandas df

In [10]: #Time based split
tscv = TimeSeriesSplit(n_splits=5)
print(tscv)

TimeSeriesSplit(max_train_size=None, n_splits=5)

In [11]: #Splitting the Dataset into Train set & Test set
for train_index, test_index in tscv.split(X):
    print("TRAIN:", train_index, "TEST:", test_index)
    X_tr, X_test = X[train_index], X[test_index]
    Y_tr, Y_test = Y[train_index], Y[test_index]

TRAIN: [ 0 1 2 ... 4997 4998 4999] TEST: [5000 5001 5002 ... 9997 9998 9999]
TRAIN: [ 0 1 2 ... 9997 9998 9999] TEST: [10000 10001 10002 ... 14997 14998 14999]
TRAIN: [ 0 1 2 ... 14997 14998 14999] TEST: [15000 15001 15002 ... 19997 19998 19999]
TRAIN: [ 0 1 2 ... 19997 19998 19999] TEST: [20000 20001 20002 ... 24997 24998 24999]
TRAIN: [ 0 1 2 ... 24997 24998 24999] TEST: [25000 25001 25002 ... 29997 29998 29999]

In [12]: # Splitting the Train dataset into Cross Validation & Train Datasets
for train_index, test_index in tscv.split(X_tr):
    print("TRAIN:", train_index, "TEST:", test_index)
    X_train, X_cv = X[train_index], X[test_index]
    Y_train, Y_cv = Y[train_index], Y[test_index]

TRAIN: [ 0 1 2 ... 4167 4168 4169] TEST: [4170 4171 4172 ... 8333 8334 8335]
TRAIN: [ 0 1 2 ... 8333 8334 8335] TEST: [ 8336 8337 8338 ... 12499 12500 12501]
TRAIN: [ 0 1 2 ... 12499 12500 12501] TEST: [12502 12503 12504 ... 16665 16666 16667]
TRAIN: [ 0 1 2 ... 16665 16666 16667] TEST: [16668 16669 16670 ... 20831 20832 20833]
TRAIN: [ 0 1 2 ... 20831 20832 20833] TEST: [20834 20835 20836 ... 24997 24998 24999]

In [13]: print(X_train.shape)
print(X_cv.shape)
print(X_test.shape)
print('-----')
print(Y_train.shape)
print(Y_cv.shape)
print(Y_test.shape)

(20834, 12)
(4166, 12)
(5000, 12)
-----
(20834,)
(4166,)
(5000,)

In [14]: # Converting X_Train, X_cv & X_test data is to Dataframe for the ease of use
X_train_data=pd.DataFrame(X_train, columns=['index', 'Id', 'ProductId', 'UserId', 'ProfileName',
      'HelpfulnessNumerator', 'HelpfulnessDenominator', 'Score', 'Time',
      'Summary', 'Text', 'CleanedText'])

X_cv_data=pd.DataFrame(X_cv, columns=['index', 'Id', 'ProductId', 'UserId', 'ProfileName',
      'HelpfulnessNumerator', 'HelpfulnessDenominator', 'Score', 'Time',
      'Summary', 'Text', 'CleanedText'])

X_test_data=pd.DataFrame(X_test, columns=['index', 'Id', 'ProductId', 'UserId', 'ProfileName',
      'HelpfulnessNumerator', 'HelpfulnessDenominator', 'Score', 'Time',
      'Summary', 'Text', 'CleanedText'])
```

Bag of Words

```
In [118]: # create the transform
vectorizer = CountVectorizer()
# tokenize and build vocab
vectorizer.fit(X_train_data['CleanedText'].values)
# Bag of Words : Train Data Set
Train_X_vector = vectorizer.transform(X_train_data['CleanedText'].values)
# summarize encoded vector
print(Train_X_vector.shape)

(20834, 19971)

In [119]: # Cross Validation Data Set
CV_X_vector = vectorizer.transform(X_cv_data['CleanedText'].values)
print(CV_X_vector.shape)

(4166, 19971)
```

```
In [120]: #Test Data Set
Test_X_vector = vectorizer.transform(X_test_data['CleanedText']).values)
print(Test_X_vector.shape)

(5000, 19971)

In [121]: print(Train_X_vector.shape)
print(CV_X_vector.shape)
print(Test_X_vector.shape)
print('-----')
print(Y_train.shape)
print(Y_cv.shape)
print(Y_test.shape)

(20834, 19971)
(4166, 19971)
(5000, 19971)
-----
(20834,)
(4166,)
(5000,)

In [122]: svd = TruncatedSVD(n_components=50)

In [123]: svd.fit(Train_X_vector)

Out[123]: TruncatedSVD(algorithm='randomized', n_components=50, n_iter=5,
    random_state=None, tol=0.0)

In [124]: TF_svd_Train=svd.transform(Train_X_vector)

In [125]: TF_svd_Train.shape

Out[125]: (20834, 50)

In [126]: TF_svd_CV=svd.transform(CV_X_vector)

In [127]: TF_svd_Test=svd.transform(Test_X_vector)

In [128]: start_time_code = time.time()
for i in range(1,30,2):
    # instantiate learning model (k = 30)
    knn = KNeighborsClassifier(n_neighbors=i, algorithm="kd_tree")

    # fitting the model on crossvalidation train
    knn.fit(TF_svd_Train, Y_train)

    # predict the response on the crossvalidation train
    pred = knn.predict(TF_svd_CV)

    # evaluate CV accuracy
    bag_acc = accuracy_score(Y_cv, pred, normalize=True) * float(100)
    print('\nCV accuracy for k = %d is %0.1f%%' % (i, bag_acc))
end time code = time.time()
print ("Running Time for code execution " + str(end_time_code - start_time_code) + " secs")

CV accuracy for k = 1 is 77.8%

CV accuracy for k = 3 is 81.2%

CV accuracy for k = 5 is 82.3%

CV accuracy for k = 7 is 82.2%

CV accuracy for k = 9 is 82.6%

CV accuracy for k = 11 is 82.6%

CV accuracy for k = 13 is 82.5%

CV accuracy for k = 15 is 82.6%

CV accuracy for k = 17 is 82.5%

CV accuracy for k = 19 is 82.6%

CV accuracy for k = 21 is 82.7%

CV accuracy for k = 23 is 82.8%

CV accuracy for k = 25 is 82.8%

CV accuracy for k = 27 is 82.7%

CV accuracy for k = 29 is 82.7%
Running Time for code execution 174.01987504959106 secs

In [129]: knn = KNeighborsClassifier(17,algorithm="kd_tree")
knn.fit(TF_svd_Train,Y_train)
pred = knn.predict(TF_svd_Test)
bag_test_acc = accuracy_score(Y_test, pred, normalize=True) * float(100)
print('\n***Test accuracy for k = 17 is %0.2f%%' % (bag_test_acc))

***Test accuracy for k = 17 is 81.94%

In [130]: Conf_matrix=confusion_matrix(Y_test, pred)

In [131]: Conf_matrix

Out[131]: array([[ 68,  865],
    [ 38, 4029]])

In [132]: plot_confusion_matrix(cm
    normalize = False,
    target_names = ['Negative', 'Positive'],
    title = "Confusion Matrix")

Confusion Matrix
True label \ Predicted label
Negative Positive
Negative 68 865
Positive 38 4029
accuracy=0.8194, misclass=0.1806
```

```
In [133]: TN,FP,FN,TP = Conf_matrix.ravel()
#confusion_matrix(Y_test, pred)
```

```
In [134]: # Sensitivity, hit rate, recall, or true positive rate
TPR = '{0:.2%}'.format( TP/(TP+FN))
# Specificity or true negative rate
TNR = '{0:.2%}'.format( TN/(TN+FP))
# Fall out or false positive rate
FPR = '{0:.2%}'.format( FP/(FP+TN))
# False negative rate
FNR = '{0:.2%}'.format( FN/(TP+FN))
# Precision or positive predictive value
PPV = '{0:.2%}'.format( TP/(TP+FP))
# Negative predictive value
NPV = '{0:.2%}'.format( TN/(TN+FN))
# Overall accuracy
ACC = '{0:.2%}'.format( (TP+TN)/(TP+FP+FN+TN))

In [135]: Recall=recall_score(Y_test, pred, average='micro')

In [136]: Precision=precision_score(Y_test, pred, average='micro')

In [137]: F1_Score=f1_score(Y_test, pred, average='weighted')

In [138]: pd.crosstab(Y_test, pred, rownames=['True'], colnames=['Predicted'], margins=True)

Out[138]:
      Predicted negative positive  All
      True
negative      68      865    933
positive      38     4029   4067
      All      106     4894   5000

In [139]: x.add_row(["Bag of Words", "kd_tree",17,TPR,TNR, '{0:.2f}'.format(Precision), '{0:.2f}'.format(Recall), '{0:.2f}'.format(F1_Score),FPR,FNR,PPV,NPV,ACC])

In [140]: print(x)

|   Model   | Algorithm | K | TPR | TNR | Precision | Recall | F1-Score | FPR | FNR | PPV | NPV | Overall Accuracy (ACC) |
| Bag of Words | kd_tree | 17 | 99.07% | 7.29% | 0.82 | 0.82 | 0.76 | 92.71% | 0.93% | 82.33% | 64.15% | 81.94% |

In [141]: print(classification_report(Y_test, pred))

              precision    recall  f1-score   support

 negative      0.64      0.07      0.13         933
 positive      0.82      0.99      0.90        4067

 avg / total      0.79      0.82      0.76        5000

In [142]: with open('Output_kd-tree.txt', 'w') as w:
w.write(str(x))
```

TF - IDF

```
In [143]: tf_transformer = TfidfVectorizer(ngram_range=(1,2))
TF_vector=tf_transformer.fit(X_train_data['CleanedText'].values)
TF_train_Vector = TF_vector.transform(X_train_data['CleanedText'].values)
TF_train_Vector.shape

Out[143]: (20834, 427963)

In [144]: TF_cv_Vector = TF_vector.transform(X_cv_data['CleanedText'].values)
TF_cv_Vector.shape

Out[144]: (4166, 427963)

In [145]: TF_test_Vector = TF_vector.transform(X_test_data['CleanedText'].values)
TF_test_Vector.shape

Out[145]: (5000, 427963)

In [146]: print(TF_train_Vector.shape)
print(TF_cv_Vector.shape)
print(TF_test_Vector.shape)
print('-----')
print(Y_train.shape)
print(Y_cv.shape)
print(Y_test.shape)

(20834, 427963)
(4166, 427963)
(5000, 427963)
-----
(20834,)
(4166,)
(5000,)

In [147]: svd = TruncatedSVD(n_components=20)

In [148]: svd.fit(TF_train_Vector)

Out[148]: TruncatedSVD(algorithm='randomized', n_components=20, n_iter=5,
random_state=None, tol=0.0)

In [149]: TF_svd_Train=svd.transform(TF_train_Vector)

In [150]: TF_svd_Train.shape

Out[150]: (20834, 20)

In [151]: TF_svd_CV=svd.transform(TF_cv_Vector)

In [152]: TF_svd_Test=svd.transform(TF_test_Vector)
```

```
In [153]: start_time_code = time.time()
for i in range(1,30,2):
    # instantiate learning model (k = 30)
    knn = KNeighborsClassifier(n_neighbors=i, algorithm="kd_tree")

    # fitting the model on crossvalidation train
    knn.fit(TF_svd_Train, Y_train)

    # predict the response on the crossvalidation train
    pred = knn.predict(TF_svd_CV)

    # evaluate CV accuracy
    tf_acc = accuracy_score(Y_cv, pred, normalize=True) * float(100)
    print('\nCV accuracy for k = %d is %0.2f%%' % (i, tf_acc))
end_time_code = time.time()
print ("Running Time for code execution " + str(end_time_code - start_time_code) + " secs")

CV accuracy for k = 1 is 78.06%

CV accuracy for k = 3 is 81.04%

CV accuracy for k = 5 is 82.72%

CV accuracy for k = 7 is 82.53%

CV accuracy for k = 9 is 82.53%

CV accuracy for k = 11 is 82.55%

CV accuracy for k = 13 is 82.93%

CV accuracy for k = 15 is 82.74%

CV accuracy for k = 17 is 82.69%

CV accuracy for k = 19 is 82.81%

CV accuracy for k = 21 is 82.74%

CV accuracy for k = 23 is 82.84%

CV accuracy for k = 25 is 82.72%

CV accuracy for k = 27 is 82.81%

CV accuracy for k = 29 is 82.79%
Running Time for code execution 46.00462007522583 secs

In [154]: knn = KNeighborsClassifier(17,algorithm="kd_tree")
knn.fit(TF_svd_Train,Y_train)
pred = knn.predict(TF_svd_Test)
tf_test_acc = accuracy_score(Y_test, pred, normalize=True) * float(100)
print('\n****Test accuracy for k = 17 is %0.2f%%' % (tf_test_acc))

****Test accuracy for k = 17 is 82.64%

In [155]: Conf_matrix=confusion_matrix(Y_test, pred)

In [156]: Conf_matrix

Out[156]: array([[ 102,  831],
 [   37, 4030]])

In [157]: plot_confusion_matrix(cm          = Conf_matrix,
                               normalize   = False,
                               target_names= ['Negative', 'Positive'],
                               title       = "Confusion Matrix")
```

Confusion Matrix

	Predicted Negative	Predicted Positive	
True Negative	102	831	
True Positive	37	4030	
			4000

accuracy=0.8264, misclass=0.1736

```
In [158]: TN,FP,FN,TP = Conf_matrix.ravel()
#confusion_matrix(Y_test, pred)

In [159]: # Sensitivity, hit rate, recall, or true positive rate
TPR = '{0:.2%}'.format( TP/(TP+FN))
# Specificity or true negative rate
TNR = '{0:.2%}'.format( TN/(TN+FP))
# Fall out or false positive rate
FPR = '{0:.2%}'.format( FP/(FP+TN))
# False negative rate
FNR = '{0:.2%}'.format( FN/(TP+FN))
# Precision or positive predictive value
PPV = '{0:.2%}'.format( TP/(TP+FP))
# Negative predictive value
NPV = '{0:.2%}'.format( TN/(TN+FN))
# Overall accuracy
ACC = '{0:.2%}'.format( (TP+TN)/(TP+FP+FN+TN))

In [160]: Recall=recall_score(Y_test, pred, average='micro')

In [161]: Precision=precision_score(Y_test, pred, average='micro')

In [162]: F1_Score=f1_score(Y_test, pred, average='weighted')

In [163]: pd.crosstab(Y_test, pred, rownames=['True'], colnames=['Predicted'], margins=True)

Out[163]:
```

	Predicted negative	positive	All
True			
negative	102	831	933
positive	37	4030	4067
All	139	4861	5000

```
In [164]: x.add_row([" TF - IDF ", "Kd_tree",17,TPR,TNR,'{0:.2f}'.format(Precision),'{0:.2f}'.format(Recall),'{0:.2f}'.format(F1_Score),FPR,FNR,PPV,NPV,ACC])

In [165]: print(classification_report(Y_test, pred))
```

	precision	recall	f1-score	support
negative	0.73	0.11	0.19	933
positive	0.83	0.99	0.90	4067
avg / total	0.81	0.83	0.77	5000

In [166]:

```
print(x)
```

	Model	Algorithm	K	TPR	TNR	Precision	Recall	F1-Score	FPR	FNR	PPV	NPV	Overall Accuracy (ACC)
Bag of Words		Kd_tree	17	99.07%	7.29%	0.82	0.82	0.76	92.71%	0.93%	82.33%	64.15%	81.94%
TF - IDF		Kd_tree	13	99.09%	10.93%	0.83	0.83	0.77	89.07%	0.91%	82.90%	73.38%	82.64%

Average W2V

In [167]: # Train your own Word2Vec model using your own text corpus for Training Dataset

```
i=0
list_of_sent=[]
for sent in X_train_data['CleanedText'].values:
    list_of_sent.append(sent.split())

# min_count = 5 considers only words that occurred atleast 5 times
Trained_model=Word2Vec(list_of_sent,min_count=5,size=50, workers=4)

w2v words = list(Trained_model.wv.vocab)
print("number of words that occurred minimum 5 times ",len(w2v_words))
print("sample words ", w2v_words[0:50])

number of words that occurred minimum 5 times 6969
sample words ['littl', 'book', 'make', 'son', 'laugh', 'loud', 'car', 'drive', 'along', 'always', 'sing', 'refrain', 'hes', 'learn', 'india', 'love', 'new', 'world', 'introduc', 'silli', 'classic', 'will', 'bet', 'still', 'abl', 'memori', 'colleg', 'rememb', 'see', 'show', 'air', 'televis', 'year', 'ago', 'child', 'sister', 'later', 'bought', 'day', 'thirti', 'someth', 'use', 'seri', 'song', 'student', 'teach', 'turn', 'whole', 'school', 'purchas']
```

In [168]: # average Word2Vec
# compute average word2vec for each review.
Trained\_vectors = []; # the avg-w2v for each sentence/review is stored in this list
for sent in list\_of\_sent: # for each review/sentence
 sent\_vec = np.zeros(50) # as word vectors are of zero length
 cnt\_words =0; # num of words with a valid vector in the sentence/review
 for word in sent: # for each word in a review/sentence
 if word in w2v\_words:
 vec = Trained\_model.wv[word]
 sent\_vec += vec
 cnt\_words += 1
 if cnt\_words != 0:
 sent\_vec /= cnt\_words
 Trained\_vectors.append(sent\_vec)
print(len(Trained\_vectors))

20834

In [169]: Train\_X=Trained\_vectors

In [170]: # Train your own Word2Vec model using your own text corpus

```
i=0
list_of_sent_cv=[]
for sent in X_cv_data['CleanedText'].values:
    list_of_sent_cv.append(sent.split())
```

In [171]: # average Word2Vec
# compute average word2vec for each review.
Cv\_vectors = []; # the avg-w2v for each sentence/review is stored in this list
for sent in list\_of\_sent\_cv: # for each review/sentence
 sent\_vec = np.zeros(50) # as word vectors are of zero length
 cnt\_words =0; # num of words with a valid vector in the sentence/review
 for word in sent: # for each word in a review/sentence
 if word in w2v\_words:
 vec = Trained\_model.wv[word]
 sent\_vec += vec
 cnt\_words += 1
 if cnt\_words != 0:
 sent\_vec /= cnt\_words
 Cv\_vectors.append(sent\_vec)
print(len(Cv\_vectors))

4166

In [172]: CV\_X=Cv\_vectors

In [173]: # Train your own Word2Vec model using your own text corpus

```
i=0
list_of_sent_test=[]
for sent in X_test_data['CleanedText'].values:
    list_of_sent_test.append(sent.split())
```

In [174]: # average Word2Vec
# compute average word2vec for each review.
test\_vectors = []; # the avg-w2v for each sentence/review is stored in this list
for sent in list\_of\_sent\_test: # for each review/sentence
 sent\_vec = np.zeros(50) # as word vectors are of zero length
 cnt\_words =0; # num of words with a valid vector in the sentence/review
 for word in sent: # for each word in a review/sentence
 if word in w2v\_words:
 vec = Trained\_model.wv[word]
 sent\_vec += vec
 cnt\_words += 1
 if cnt\_words != 0:
 sent\_vec /= cnt\_words
 test\_vectors.append(sent\_vec)
print(len(test\_vectors))

5000

In [175]: Test\_X=test\_vectors

In [176]: print (len(Train\_X))
print(len(CV\_X))
print(len(Test\_X))
print('-----')
print (Y\_train.shape)
print(Y\_cv.shape)
print(Y\_test.shape)

20834

4166

5000

-----
(20834,)
(4166,)
(5000,)

```
In [177]: start_time_code = time.time()
for i in range(1,30,2):
    # instantiate learning model (k = 30)
    knn = KNeighborsClassifier(n_neighbors=i, algorithm="kd_tree")

    # fitting the model on crossvalidation train
    knn.fit(Train_X, Y_train)

    # predict the response on the crossvalidation train
    pred = knn.predict(CV_X)

    # evaluate CV accuracy
    train_acc = accuracy_score(Y_cv, pred, normalize=True) * float(100)
    print('\nCV accuracy for k = %d is %0.2f%%' % (i, train_acc))
end_time_code = time.time()
print ("Running Time for code execution " + str(end_time_code - start_time_code) + " secs")

CV accuracy for k = 1 is 79.28%

CV accuracy for k = 3 is 82.45%

CV accuracy for k = 5 is 83.39%

CV accuracy for k = 7 is 83.27%

CV accuracy for k = 9 is 83.22%

CV accuracy for k = 11 is 83.44%

CV accuracy for k = 13 is 83.34%

CV accuracy for k = 15 is 83.41%

CV accuracy for k = 17 is 83.39%

CV accuracy for k = 19 is 83.27%

CV accuracy for k = 21 is 83.34%

CV accuracy for k = 23 is 83.37%

CV accuracy for k = 25 is 83.29%

CV accuracy for k = 27 is 83.29%

CV accuracy for k = 29 is 83.27%
Running Time for code execution 153.1801438331604 secs
```

```
In [178]: # Since the train accuracy for k is the same from k=3 to k=29. so the smallest k at maximum accuracy is the best i.e. k=3
knn = KNeighborsClassifier(11,algorithm="kd_tree")
knn.fit(Train_X,Y_train)
pred = knn.predict(Test_X)
test_acc = accuracy_score(Y_test, pred, normalize=True) * float(100)
print('\n****Test accuracy for k = 11 is %0.2f%%' % (test_acc))


****Test accuracy for k = 11 is 83.30%
```

```
In [179]: Conf_matrix=confusion_matrix(Y_test, pred)
```

```
In [180]: Conf_matrix
```

```
Out[180]: array([[ 169,   764],
                 [   71, 3996]])
```

```
In [181]: plot_confusion_matrix(cm          = Conf_matrix,
                               normalize   = False,
                               target_names= ['Negative', 'Positive'],
                               title       = "Confusion Matrix")
```



```
In [182]: TN,FP,FN,TP = Conf_matrix.ravel()
#confusion_matrix(Y_test, pred)
```

```
In [183]: # Sensitivity, hit rate, recall, or true positive rate
TPR = '{0:.2%}'.format( TP/(TP+FN))
# Specificity or true negative rate
TNR = '{0:.2%}'.format( TN/(TN+FP))
# Fall out or false positive rate
FPR = '{0:.2%}'.format( FP/(FP+TN))
# False negative rate
FNR = '{0:.2%}'.format( FN/(TP+FN))
# Precision or positive predictive value
PPV = '{0:.2%}'.format( TP/(TP+FP))
# Negative predictive value
NPV = '{0:.2%}'.format( TN/(TN+FN))
# Overall accuracy
ACC = '{0:.2%}'.format( (TP+TN)/(TP+FP+FN+TN))
```

```
In [184]: Recall=recall_score(Y_test, pred, average='micro')
```

```
In [185]: Precision=precision_score(Y_test, pred, average='micro')
```

```
In [186]: F1_Score=f1_score(Y_test, pred, average='weighted')
```

```
In [187]: pd.crosstab(Y_test, pred, rownames=['True'], colnames=['Predicted'], margins=True)
```

```
Out[187]:
```

	Predicted negative	positive	All
True			
negative	169	764	933
positive	71	3996	4067
All	240	4760	5000

```
In [188]: x.add_row([" Average W2V ", "Kd_tree",11,TPR,TNR,'{0:.2%}'.format(Precision),'{0:.2f}'.format(Recall),'{0:.2f}'.format(F1_Score),FPR,FNR,PPV,NPV,ACC])
```

```
In [189]: print(classification_report(Y_test, pred))
```

	precision	recall	f1-score	support
negative	0.70	0.18	0.29	933
positive	0.84	0.98	0.91	4067
avg / total	0.81	0.83	0.79	5000

In [190]:

	Model	Algorithm	K	TPR	TNR	Precision	Recall	F1-Score	FPR	FNR	PPV	NPV	Overall Accuracy (ACC)			
	Bag Of Words	kd_tree	17	99.07%	7.29%	0.82	0.82	0.76	92.71%	0.93%	82.33%	64.15%	81.94%			
	TF - IDF	Kd_tree	13	99.09%	10.93%	0.83	0.83	0.77	89.07%	0.91%	82.90%	73.38%	82.64%			
	Average W2V	Kd_tree	11	98.25%	18.11%	0.83	0.83	0.79	81.89%	1.75%	83.95%	70.42%	83.30%			

TF-IDF weighted Word2Vec

In [53]:

```
# TF-IDF weighted Word2Vec
tfidf_feat = TF_vector.get_feature_names() # tfidf words/col-names
```

In [54]:

```
# final_tf_idf is the sparse matrix with row= sentence, col=word and cell_val = tfidf
start_time_code = time.time()
tfidf_sent_vectors = []; # the tfidf-w2v for each sentence/review is stored in this list
row=0;
for sent in list_of_sent: # for each review/sentence
    sent_vec = np.zeros(50) # as word vectors are of zero length
    weight_sum =0; # num of words with a valid vector in the sentence/review
    for word in sent: # for each word in a review/sentence
        if word in w2v_words:
            vec = Trained_model.wv[word]
            # obtain the tf_idfidf of a word in a sentence/review
            tf_idf = TF_train_Vector[row, tfidf_feat.index(word)]
            sent_vec += (vec * tf_idf)
            weight_sum += tf_idf
        if weight_sum != 0:
            sent_vec /= weight_sum
            tfidf_sent_vectors.append(sent_vec)
            row += 1
end_time_code = time.time()
print ("Running Time for code execution " + str(end_time_code - start_time_code) + " secs")
```

Running Time for code execution 13474.992048978806 secs

In [55]:

```
Train_X=tfidf_sent_vectors
```

In [56]:

```
df = pd.DataFrame.from_records(Train_X)
df.to_csv('TF-IDF_W2V_Train_kd.csv',index=None)
```

In [191]:

```
Train_X = pd.read_csv("TF-IDF_W2V_Train_kd.csv")
```

In [57]:

```
# final_tf_idf is the sparse matrix with row= sentence, col=word and cell_val = tfidf
start_time_code = time.time()
tfidf_sent_vectors = []; # the tfidf-w2v for each sentence/review is stored in this list
row=0;
for sent in list_of_sent_cv: # for each review/sentence
    sent_vec = np.zeros(50) # as word vectors are of zero length
    weight_sum =0; # num of words with a valid vector in the sentence/review
    for word in sent: # for each word in a review/sentence
        if word in w2v_words:
            vec = Trained_model.wv[word]
            # obtain the tf_idfidf of a word in a sentence/review
            tf_idf = TF_cv_Vector[row, tfidf_feat.index(word)]
            sent_vec += (vec * tf_idf)
            weight_sum += tf_idf
        if weight_sum != 0:
            sent_vec /= weight_sum
            tfidf_sent_vectors.append(sent_vec)
            row += 1
end_time_code = time.time()
print ("Running Time for code execution " + str(end_time_code - start_time_code) + " secs")
```

Running Time for code execution 2435.09094786644 secs

In [58]:

```
CV_X=tfidf_sent_vectors
```

In [59]:

```
df = pd.DataFrame.from_records(CV_X)
df.to_csv('TF-IDF_W2V_CV_Kd.csv',index=None)
```

In [192]:

```
CV_X = pd.read_csv("TF-IDF_W2V_CV_Kd.csv")
```

In [60]:

```
# final_tf_idf is the sparse matrix with row= sentence, col=word and cell_val = tfidf
start_time_code = time.time()
tfidf_sent_vectors = []; # the tfidf-w2v for each sentence/review is stored in this list
row=0;
for sent in list_of_sent_test: # for each review/sentence
    sent_vec = np.zeros(50) # as word vectors are of zero length
    weight_sum =0; # num of words with a valid vector in the sentence/review
    for word in sent: # for each word in a review/sentence
        if word in w2v_words:
            vec = Trained_model.wv[word]
            # obtain the tf_idfidf of a word in a sentence/review
            tf_idf = TF_test_Vector[row, tfidf_feat.index(word)]
            sent_vec += (vec * tf_idf)
            weight_sum += tf_idf
        if weight_sum != 0:
            sent_vec /= weight_sum
            tfidf_sent_vectors.append(sent_vec)
            row += 1
end_time_code = time.time()
print ("Running Time for code execution " + str(end_time_code - start_time_code) + " secs")
```

Running Time for code execution 2871.7937688827515 secs

In [61]:

```
Test_X=tfidf_sent_vectors
```

In [62]:

```
df = pd.DataFrame.from_records(Test_X)
df.to_csv('TF-IDF_W2V_Test_Kd.csv',index=None)
```

In [193]:

```
Test_X = pd.read_csv("TF-IDF_W2V_Test_Kd.csv")
```

```
In [194]: start_time_code = time.time()
for i in range(1,30,2):
    # instantiate learning model (k = 30)
    knn = KNeighborsClassifier(n_neighbors=i, algorithm="kd_tree")

    # fitting the model on crossvalidation train
    knn.fit(Train_X, Y_train)

    # predict the response on the crossvalidation train
    pred = knn.predict(CV_X)

    # evaluate CV accuracy
    train_acc = accuracy_score(Y_cv, pred, normalize=True) * float(100)
    print('\nCV accuracy for k = %d is %0.2f%%' % (i, train_acc))
end_time_code = time.time()
print ("Running Time for code execution " + str(end_time_code - start_time_code) + " secs")

CV accuracy for k = 1 is 77.87%

CV accuracy for k = 3 is 81.06%

CV accuracy for k = 5 is 82.53%

CV accuracy for k = 7 is 82.67%

CV accuracy for k = 9 is 83.37%

CV accuracy for k = 11 is 83.15%

CV accuracy for k = 13 is 83.08%

CV accuracy for k = 15 is 82.89%

CV accuracy for k = 17 is 82.89%

CV accuracy for k = 19 is 82.93%

CV accuracy for k = 21 is 82.89%

CV accuracy for k = 23 is 82.93%

CV accuracy for k = 25 is 82.91%

CV accuracy for k = 27 is 83.03%

CV accuracy for k = 29 is 82.86%
Running Time for code execution 135.67332410812378 secs
```

```
In [195]: # Since the train accuracy for k is the same from k=3 to k=29. so the smallest k at maximum accuracy is the best i.e. k=3
knn = KNeighborsClassifier(23,algorithm="kd_tree")
knn.fit(Train_X,Y_train)
pred = knn.predict(Test_X)
test_acc = accuracy_score(Y_test, pred, normalize=True) * float(100)
print('\n***Test accuracy for k = 23 is %0.2f%%' % (test_acc))


***Test accuracy for k = 23 is 82.84%
```

```
In [196]: Conf_matrix=confusion_matrix(Y_test, pred)
```

```
In [197]: Conf_matrix
```

```
Out[197]: array([[ 102,  831],
                [  27, 4040]])
```

```
In [198]: plot_confusion_matrix(cm          = Conf_matrix,
                                normalize   = False,
                                target_names = ['Negative', 'Positive'],
                                title       = "Confusion Matrix")
```



```
In [199]: TN,FP,FN,TP = Conf_matrix.ravel()
#confusion_matrix(Y_test, pred)
```

```
In [200]: # Sensitivity, hit rate, recall, or true positive rate
TPR = '{0:.2%}'.format( TP/(TP+FN))
# Specificity or true negative rate
TNR = '{0:.2%}'.format( TN/(TN+FP))
# Fall out or false positive rate
FPR = '{0:.2%}'.format( FP/(FP+TN))
# False negative rate
FNR = '{0:.2%}'.format( FN/(TP+FN))
# Precision or positive predictive value
PPV = '{0:.2%}'.format( TP/(TP+FP))
# Negative predictive value
NPV = '{0:.2%}'.format( TN/(TN+FN))
# Overall accuracy
ACC = '{0:.2%}'.format( (TP+TN)/(TP+FP+FN+TN))
```

```
In [201]: Recall=recall_score(Y_test, pred, average='micro')
```

```
In [202]: Precision=precision_score(Y_test, pred, average='micro')
```

```
In [203]: F1_Score=f1_score(Y_test, pred, average='weighted')
```

```
In [204]: pd.crosstab(Y_test, pred, rownames=['True'], colnames=['Predicted'], margins=True)
```

```
Out[204]:
```

	Predicted negative	positive	All
True			
negative	102	831	933
positive	27	4040	4067
All	129	4871	5000

```
In [205]: x.add_row([" TF - IDF Weighted W2V ", "Kd_tree", 23,TPR,TNR, '{0:.2f}'.format(Precision), '{0:.2f}'.format(Recall), '{0:.2f}'.format(F1_Score), FPR,FNR,PPV,NPV,ACC])
```

```
In [206]: print(classification_report(Y_test, pred))
```

	precision	recall	f1-score	support
negative	0.79	0.11	0.19	933
positive	0.83	0.99	0.90	4067
avg / total	0.82	0.83	0.77	5000



```
In [299]: print(x)
```

	Model	Algorithm	K	TPR	TNR	Precision	Recall	F1-Score	FPR	FNR	PPV	NPV	Overall Accuracy (ACC)
	Bag of Words	kd_tree	17	99.07%	7.29%	0.82	0.82	0.76	92.71%	0.93%	82.33%	64.15%	81.94%
	TF - IDF	Kd_tree	13	99.09%	10.93%	0.83	0.83	0.77	89.07%	0.91%	82.90%	73.38%	82.64%
	Average W2V	Kd_tree	11	98.25%	18.11%	0.83	0.83	0.79	81.89%	1.75%	83.95%	70.42%	83.30%
TF - IDF Weighted W2V	Bag of Words	Kd_tree	23	99.34%	10.93%	0.83	0.83	0.77	89.07%	0.66%	82.94%	79.07%	82.84%
	Bag of Words	Brute	15	99.93%	1.87%	0.82	0.82	0.74	98.13%	0.07%	81.68%	85.29%	81.70%
	TF - IDF	Brute	11	99.04%	12.27%	0.83	0.83	0.78	87.73%	0.96%	83.18%	74.51%	82.91%
	Average W2V	Brute	17	98.36%	20.85%	0.84	0.84	0.80	79.15%	1.64%	84.48%	74.42%	83.96%
	TF-IDF Weighted W2V	Brute	17	98.88%	2.07%	0.81	0.81	0.73	97.93%	1.12%	81.56%	29.63%	80.88%

```
In [300]: with open('Output.txt', 'w') as w:
          w.write(str(x))
```

Conclusion

- 1. Of all the Word to vector conversion models, TF - IDF got more True Positive Rate and True Negative Rate in 'Brute' and 'Kd-tree'